

内核调试实验

实验一：系统是如何从start_kernel开始一步步进入启动用户空间第一个程序的（提示）：

watch system_state ramdisk_execute_command

SYSTEM_BOOTING

```
(gdb) p system_state  
$1 = SYSTEM_BOOTING
```

内核进入c语言阶段，会执行start_kernel 576 初始化内核，首先完成一些核心环境的初始化

```
Breakpoint 2, start_kernel () at init/main.c:576  
576 {  
(gdb) list  
571 {  
572     rest_init();  
573 }  
574  
575 asmlinkage __visible void __init start_kernel(void)  
576 {  
577     char *command_line;  
578     char *after_dashes;  
579  
580     set_task_stack_end_magic(&init_task);  
(gdb) c  
Continuing.
```

set_task_stack_end_magic(人工0号进程) kernel/fork.c:786

```
Breakpoint 3, set_task_stack_end_magic (tsk=0xffffffff82613780 <init_task>)  
at kernel/fork.c:835  
835 {  
(gdb) list  
830     *dst = *src;  
831     return 0;  
832 }  
833  
834 void set_task_stack_end_magic(struct task_struct *tsk)  
835 {  
836     unsigned long *stackend;  
837  
838     stackend = end_of_stack(tsk);  
839     *stackend = STACK_END_MAGIC; /* for overflow detection */  
(gdb)  
840 }
```

rdinit_setup 352 如果bootargs设置了rdinit，那么内核在启动阶段会解析并赋给ramdisk_execute_command

```

Breakpoint 1, rdinit_setup (str=0xffff88803fee07d5 "/helloworld")
  at init/main.c:352
352     {
(gdb) l
347         return 1;
348     }
349     __setup("init=", init_setup);
350
351     static int __init rdinit_setup(char *str)
352     {
353         unsigned int i;
354
355         ramdisk_execute_command = str;
356         /* See "auto" comment in init_setup */
Hardware watchpoint 14: ramdisk_execute_command

Old value = 0x0 <fixed_percpu_data>
New value = 0xffff88803fee07d5 "/helloworld"
0xffffffff82876689 in rdinit_setup (str=0xffff88803fee07d5 "/helloworld") at ini
t/main.c:355
355         ramdisk_execute_command = str;

```

sched_init kernel/sched/core.c:6376 进程调度相关初始化

```

Breakpoint 4, sched_init () at kernel/sched/core.c:6376
6376     {
(gdb) l
6371
6372     DECLARE_PER_CPU(cpumask_var_t, load_balance_mask);
6373     DECLARE_PER_CPU(cpumask_var_t, select_idle_mask);
6374
6375     void __init sched_init(void)
6376     {
6377         unsigned long alloc_size = 0, ptr;
6378         int i;
6379
6380         wait_bit_init();

```

rest_init init/main.c:407 完成最后的初始化工作

```

Breakpoint 5, rest_init () at init/main.c:407
407     {
(gdb) l
402     */
403
404     static __initdata DECLARE_COMPLETION(kthreadd_done);
405
406     noinline void __ref rest_init(void)
407     {
408         struct task_struct *tsk;
409         int pid;
410
411         rcu_scheduler_starting();
[ 0.000000] Linux version 5.3.0 (amos@ubuntu) (gcc version 7.4.0 (Ubuntu 7.4.
0-1ubuntu1~18.04.1)) #1 SMP Sat Feb 15 19:31:58 +08 2020
[ 0.000000] Command line: root=/dev/sda rdinit=/helloworld nokaslr console=tt
ys0

```

kernel_init(->1号进程, 先创建, 但是要等待2号进程已创建完成, 化身为用户进程祖先) 1107

```

Breakpoint 6, kernel_init (unused=0x0 <fixed_percpu_data>) at init/main.c:1107
1107     {
(gdb) l
1102     {
1103         free_initmem_default(POISON_FREE_INITMEM);
1104     }
1105
1106     static int __ref kernel_init(void *unused)
1107     {
1108         int ret;
1109
1110         kernel_init_freeable();
1111         /* need to finish all async __init code before freeing the memor
v */

```

kernel_init_freeable 1160 -> do_basic_setup 1001 -> driver_init 注册内核驱动模块

```

Breakpoint 7, kernel_init_freeable () at init/main.c:1160
1160     {
(gdb) l
1155         panic("No working init found. Try passing init= option to kerne
l. "
1156         "See Linux Documentation/admin-guide/init.rst for guidance
.");
1157     }
1158
1159     static noinline void __init kernel_init_freeable(void)
1160     {
1161         /*
1162         * Wait until kthreadd is all set-up.
1163         */
1164         wait_for_completion(&kthreadd_done);

```

SYSTEM_SCHEDULING此时系统已经初始化完毕, 但是必要的内核线程还没有运行起来, 此时0号和1号进程卡在这里等待2号进程以及它的一系列必要子进程创建完成

Hardware watchpoint 13: system_state

```
Old value = SYSTEM_BOOTING
New value = SYSTEM_SCHEDULING
rest_init () at init/main.c:443
443         complete(&kthreadd_done);
(gdb) l
438         * CONFIG_PREEMPT_VOLUNTARY=y the init task might have scheduled
439         * already, but it's stuck on the kthreadd_done completion.
440         */
441         system_state = SYSTEM_SCHEDULING;
442
443         complete(&kthreadd_done);
444
```

kthreadd(->2号进程, 所有内核线程的祖先) kernel/kthread.c:215

```
Breakpoint 8, kthread (_create=0xfffff88803e544c40) at kernel/kthread.c:215
215     {
(gdb) l
210         __kthread_parkme(to_kthread(current));
211     }
212     EXPORT_SYMBOL_GPL(kthread_parkme);
213
214     static int kthread(void *_create)
215     {
216         /* Copy data: it's on kthread's stack */
217         struct kthread_create_info *create = _create;
218         int (*threadfn)(void *data) = create->threadfn;
219         void *data = create->data;
```

kgdboc调试断在这里。

内核线程初始化之后, 1号进程可以进行do_basic_setup如driver_init等 完成设备、驱动等初始化

```
Breakpoint 11, kernel_init_freeable () at init/main.c:1192
1192         do_basic_setup();
(gdb) l
1187
1188         page_alloc_init_late();
1189         /* Initialize page ext after all struct pages are initialized.
1190         /
1190         page_ext_init();
1191
1192         do_basic_setup();
```

cpu_startup_entry(0号进程) kernel/sched/idle.c:350 在完成全部的启动后, 进入idle循环, 化身为idle进程

```

Breakpoint 10, cpu_startup_entry (state=CPUHP_ONLINE)
  at kernel/sched/idle.c:350
350     {
(gdb) l
345         preempt_enable();
346     }
347     EXPORT_SYMBOL_GPL(play_idle);
348
349     void cpu_startup_entry(enum cpuhp_state state)
350     {
351         arch_cpu_idle_prepare();
352         cpuhp_online_idle(state);
353         while (1)
354             do_idle();

```

SYSTEM_RUNNING此时系统才算运行起来

```

Hardware watchpoint 13: system_state

Old value = SYSTEM_SCHEDULING
New value = SYSTEM_RUNNING
kernel_init (unused=<optimized out>) at init/main.c:1124
1124     numa_default_policy();
(gdb) l
1119         * to finalize PTI.
1120         */
1121     pti_finalize();
1122
1123     system_state = SYSTEM_RUNNING;
1124     numa_default_policy();
1125
1126     rcu_end_inkernel_boot();
1127
1128     if (ramdisk_execute_command) {

```

run_init_process 1045 kernel_init作为1号进程化身为用户空间祖先

```

Breakpoint 9, run_init_process (init_filename=0xffff88803fee07d5 "/helloworld")
  at init/main.c:1045
1045  {
(gdb) l
1040      for (fn = __initcall_start; fn < __initcall0_start; fn++)
1041          do_one_initcall(initcall_from_entry(fn));
1042  }
1043
1044  static int run_init_process(const char *init_filename)
1045  {
1046      argv_init[0] = init_filename;
1047      pr_info("Run %s as init process\n", init_filename);
1048      return do_execve(getname_kernel(init_filename),
1049                      (const char __user *const __user *)argv_init,
(gdb)
1050                      (const char __user *const __user *)envp_init);
1051  }

```

进入系统

```

[ 5.832170] Run /helloworld as init process
Hello World
This is an entry
Author:your own name

```

实验二（提示）：在qemu里调试模块

1.将模块添加入到制作的busybox.img中的/lib目录下

```
>> sudo mount -o loop ~/kDebug/busybox.img /mnt/disk
```

```
>> sudo cp ~/kDebug/hello.ko /mnt/disk/lib
```

```
>> sudo umount /mnt/disk
```

2.sudo qemu-system-x86_64 -kernel bzImage -initrd initrd.img-5.3.0 -append "root=/dev/sda nokaslr" -boot c -hda busybox.img -k en-us -m 1024 -serial tcp::4321,server

不加kgdbwait使能参数

开发机gdb连接

3.启动qemu之后insmod hello

```

/ # cd lib
/lib # insmod hello.ko test=1111
[ 117.747978] hello: loading out-of-tree module taints kernel.
[ 117.748560] hello: module verification failed: signature and
missing - tainting kernel
[ 117.755677] Hello guoqi test=1111

```

4.使能kgdb, echo ttyS0 > /sys/module/kgdboc/parameters/kgdboc 注册kgdb

然后触发断点 `echo g > /proc/sysrq-trigger`

```
/ # echo ttyS0 > /sys/module/kgdboc/parameters/kgdboc
[ 52.835810] KGDB: Registered I/O driver kgdboc
/ # echo g > /proc/sysrq-trigger
[ 66.225220] sysrq: DEBUG

Entering kdb (current=0xffff88803dae2b80, pid 264) on processor 0 due to Keyboard Entry
```

这时开发机gdb再连接，会成功断在kgdb_breakpoint()这里

5.开发机lx-symbols把hello.ko模块中的符号加入进来，这时可设置hello里的函数断点、watch变量，并print 变量

```
(gdb) lx-symbols
loading vmlinux
scanning for modules in /home/amos/kDebug
loading @0xfffffffffc0035000: /home/amos/kDebug/hello.ko
(gdb) watch test
Hardware watchpoint 1: test
(gdb) p test
$1 = 1111
(gdb) c
Continuing.
```

6.按c把主动权回到调试机

reference:

<https://01.org/linuxgraphics/gfx-docs/drm/dev-tools/gdb-kernel-debugging.html>